

QA Engineer Assessment — Response

GitHub Repository: [Automated-Tests-VoiceScript](#)

Candidate: Yusack Manuel Date: 2026-06-01 Case: Moonlight Plaza Associates v. Cancun Farms

Task 1 — Identify Issues (15 Found)

Category A: Transcription Errors (ASR Misrecognitions)

#	Issue	Location	Why It Matters
1	" scientists " instead of "clients"	Raw ASR: we represent scientists, Moonlight Plaza Association	Completely changes the meaning — the firm represents legal clients, not scientists. Critical legal error.
2	" Moonlite " instead of "Moonlight"	Raw ASR word: Moonlite	Misspells the plaintiff's name. Undermines credibility of the transcript.
3	" councel " instead of "counsel"	Raw ASR word: councel.	Misspells a key legal term.
4	" Richman & Lavine " instead of "Richmond & Levine"	Processed transcript: Richman & Lavine PC vs. raw ASR: Richmond and Levine	The AI processing introduced a new error. The raw ASR was correct ("Richmond and Levine") but the AI "corrected" it to the wrong name.
5	" Cancun Farms " / " Cancun Forms " instead of "Cannon Farms"	Processed transcript: multiple instances of Cancun Farms, Cancun Forms 2	Consistently misnames the defendant party. "Cancun" is a city in Mexico — not a farm company.
6	" No Further Action leather " instead of "No Further Action letter"	Processed transcript	A "No Further Action letter" is a specific regulatory document. "Leather" changes the meaning entirely.
7	" Professional Geologists " instead of "Professional Geologist"	Processed transcript (suffix PG)	PG = Professional Geologist (singular). The plural form is incorrect for a single-person suffix.
8	" Miss Jacob " instead of "Ms. Jacob"	Raw ASR word: Miss	ASR misrecognized "Ms." as "Miss." Minor but shows ASR quality issues with honorifics.

Category B: Speaker Attribution Issues

#	Issue	Location	Why It Matters
9	Speaker labels swapped in metadata	metadata_corrupted.json: Speaker B (Chris Jacob, PLAINTIFF ATTORNEY) has real_time_asr_label: "THE WITNESS". Speaker A (Terry/Jerry	Speaker B is labeled as PLAINTIFF ATTORNEY but Chris Jacob is the witness (project manager at Cannon Farms).

		Sellingman, WITNESS) has real_time_asr_label: "MS SELIGMAN"	Speaker A is labeled WITNESS but Jerry Sellingman is the attorney. Roles are completely reversed.
10	"MS. ATRIANO" appears in processed transcript	Multiple lines in processed transcript	No speaker named "Ms. Atriano" exists in the metadata. This is an unknown/unmapped speaker label — could be the defendant's attorney objecting, but the system has no record of this person.
11	Witness answers attributed to wrong speaker	Processed transcript: WITNESS: I am a senior environmental project manager. And what are your duties...	The question "And what are your duties..." is the attorney's question, not the witness's answer. It's incorrectly lumped under the WITNESS label.
12	"WITNESS" label used inconsistently	Processed transcript uses both WITNESS and MS SELIGMAN for the same person at different points	Inconsistent label normalization — the AI sometimes uses the role-based label and sometimes the name-based label.
13	Name mismatch in metadata Speaker A	first_name: "Terry", last_name: "Sellingman", but full_name: "Jerry Sellingman"	First name and full name disagree. The ASR correctly identifies "Jerry Sellingman" but the first_name field says "Terry."

Category C: Timestamp Inconsistencies

#	Issue	Location	Why It Matters
14	Word timestamps precede utterance start time	Raw ASR: Utterance start=17920, but first word "Good" starts at 16000	A word cannot start before its containing utterance. This indicates a data integrity bug in the ASR pipeline.
15	Duplicate word timestamps	Raw ASR: Words "I" and "am" both have start: 4560; words "a" and "little" both have start: 38240; words "you" and "do" both have start: 174080	Two distinct words cannot occupy the exact same millisecond. This breaks any time-based indexing or playback sync.
16	created_at timestamp is invalid	Metadata: "created_at": "2025-10-09T26:61:00Z"	Hour 26 and minute 61 are both invalid. This would crash any parser expecting valid ISO 8601.
17	Audio/transcript duration mismatch	Metadata: audio_duration_ms: 630000 (10:30) vs. transcript_duration_ms: 605000 (10:05)	25-second gap is unaccounted for. Could

			mean missing content or a processing error.
--	--	--	---

Category D: Formatting / Structural Problems

#	Issue	Location	Why It Matters
18	Duplicate lines in processed transcript	THE REPORTER: Absolutely. appears twice; THE REPORTER: Good morning... appears in both timestamped and non-timestamped sections	Duplication inflates the transcript and confuses readers. Indicates a bug in the AI deduplication logic.
19	"error" field is a string, not null	Metadata: "error": "null"	Should be null (JSON null) or omitted. A string "null" would pass a truthy check in many languages, masking the absence of an error.
20	Transcript ends mid-sentence	Processed transcript ends with MS SELIGMAN: And going.	An incomplete final utterance indicates truncated processing or a stream cutoff bug.

Task 2 — Test Cases (10 Designed)

TC-1: ASR Word Accuracy — Known Entity Names

Field	Value
Description	Verify that known entity names (party names, firm names) are transcribed correctly
Input	Raw ASR transcript containing "Moonlight Plaza Association", "Cannon Farms", "Richmond & Levine"
Expected Output	All entity names match the metadata parties list exactly
Failure Condition	Any entity name has Levenshtein distance > 0 from the canonical name in metadata

TC-2: Speaker Role Consistency

Field	Value
Description	Verify that speaker roles in metadata are internally consistent (first_name + last_name = full_name)
Input	metadata_corrupted.json speakers array
Expected Output	full_name === first_name + " " + last_name for all speakers

Failure Condition	Any speaker has mismatched name components
--------------------------	--

TC-3: Timestamp Monotonicity Within Utterances

Field	Value
Description	Verify that word timestamps within an utterance are monotonically non-decreasing
Input	Raw ASR utterance with word-level timestamps
Expected Output	For each utterance, <code>words[i].start <= words[i+1].start</code> for all <code>i</code>
Failure Condition	Any word has a timestamp earlier than the previous word

TC-4: Utterance-Word Timestamp Containment

Field	Value
Description	Verify that all word timestamps fall within their parent utterance's time range
Input	Raw ASR utterance with start time and word array
Expected Output	<code>word.start >= utterance.start</code> for all words
Failure Condition	Any word starts before its utterance

TC-5: Duplicate Timestamp Detection

Field	Value
Description	Detect words with identical timestamps within the same utterance
Input	Raw ASR utterance word array
Expected Output	No two consecutive words share the exact same timestamp
Failure Condition	Any pair of words has <code>word[i].start === word[i+1].start</code>

TC-6: ISO 8601 Date Validation

Field	Value
Description	Verify all date/time fields in metadata are valid ISO 8601
Input	<code>metadata_corrupted.json</code> — <code>created_at</code> field
Expected Output	<code>Date.parse(created_at)</code> returns a valid date; hours 0-23, minutes 0-59
Failure Condition	<code>isNaN(Date.parse(created_at))</code> or extracted hour > 23 or minute > 59

TC-7: Audio-Transcript Duration Reconciliation

Field	Value
Description	Verify transcript duration does not exceed audio duration
Input	Metadata audio_duration_ms and transcript_duration_ms
Expected Output	transcript_duration_ms <= audio_duration_ms
Failure Condition	transcript_duration_ms > audio_duration_ms or gap > threshold (e.g., 5000ms)

TC-8: Speaker Label Mapping Completeness

Field	Value
Description	Verify every speaker label in the transcript has a corresponding entry in metadata
Input	Set of unique speaker labels from transcript + metadata speakers array
Expected Output	Every transcript speaker label matches a post_asr_label or real_time_asr_label in metadata
Failure Condition	Any transcript speaker label (e.g., "MS. ATRIANO") has no metadata mapping

TC-9: Duplicate Line Detection in Processed Transcript

Field	Value
Description	Detect consecutive duplicate lines in the processed transcript
Input	Processed transcript text file
Expected Output	No two consecutive non-empty lines are identical
Failure Condition	Any line N === line N+1 (after trimming whitespace)

TC-10: Transcript Completeness Check

Field	Value
Description	Verify the transcript does not end with a sentence fragment
Input	Last line of processed transcript
Expected Output	Last utterance ends with terminal punctuation (. ! ?) or is a complete sentence
Failure Condition	Last line ends with a conjunction or incomplete phrase (e.g., "And going.")

Task 3 — API Testing Plan

Endpoint: `POST /transcripts/process`

Purpose

Submits raw ASR output + metadata for processing; returns a processed transcript.

Test Strategy

Happy Path:

- Submit valid raw ASR JSON + valid metadata → expect 200 OK with processed transcript
- Verify response contains all utterances from input
- Verify speaker labels are normalized consistently

Validation Rules:

Rule	Test
Required fields	Omit utterances → expect 400 Bad Request
Required fields	Omit speakers in metadata → expect 400
Timestamp format	Submit invalid timestamps (negative, non-numeric) → expect 400
Speaker mapping	Submit transcript with unmapped speaker labels → expect 400 or warning
Empty payload	Submit {} → expect 400
Content-Type	Submit without application/json → expect 415

Edge Cases:

- Single-word utterance
- Very long transcript (1000+ utterances) — test performance/timeouts
- Overlapping speaker timestamps (two speakers at same time)
- Zero-duration audio (empty file)
- Unicode/special characters in transcript text
- Null/missing word-level timestamps
- Duplicate speaker labels in metadata

Response Validation:

- Response schema matches expected format
- All input utterances are represented in output
- No hallucinated content (output doesn't add words not in input)
- Speaker labels are consistently normalized

Endpoint: GET /transcripts/:id

Purpose

Retrieves a previously processed transcript by ID.

Test Strategy

Happy Path:

- Submit valid transcript ID → expect 200 OK with transcript data
- Verify returned data matches what was stored

Validation Rules:

Rule	Test
------	------

Valid ID format	Submit malformed ID (e.g., "abc-123!") → expect 400
Missing ID	Submit empty path → expect 400 or 404
Non-existent ID	Submit valid-format but non-existent ID → expect 404

Edge Cases:

- ID with SQL injection attempts (`1; DROP TABLE`)
- ID with XSS attempts (`<script>alert(1)</script>`)
- Extremely long ID string
- Concurrent requests for same ID — verify consistent response
- Transcript still processing (if async) → expect `202 Accepted` or appropriate status
- Deleted/archived transcript → expect `404` or `410 Gone`

Response Validation:

- Include `Content-Type: application/json`
- Include transcript metadata (created_at, duration, speaker count)
- Include word-level timestamps if requested
- Consistent response across repeated calls (idempotency)

Task 4 — Workflow Testing

State Machine: `NEW` → `ASSIGNED` → `TRANSCRIBED` → `REVIEWED` → `COMPLETED`

Valid Transition Tests

Test	From	Action	To	Expected
WT-1	NEW	Assign to user	ASSIGNED	200, status = ASSIGNED
WT-2	ASSIGNED	Submit transcript	TRANSCRIBED	200, status = TRANSCRIBED
WT-3	TRANSCRIBED	Submit review	REVIEWED	200, status = REVIEWED
WT-4	REVIEWED	Mark complete	COMPLETED	200, status = COMPLETED

Invalid Transition Tests

Test	From	Action	To	Expected
WT-5	NEW	Submit transcript	—	409 Conflict — must be ASSIGNED first
WT-6	NEW	Mark complete	—	409 Conflict — cannot skip steps
WT-7	ASSIGNED	Mark complete	—	409 Conflict — must be TRANSCRIBED first
WT-8	TRANSCRIBED	Mark complete	—	409 Conflict — must be REVIEWED first
WT-9	COMPLETED	Re-assign	—	409 Conflict — terminal state
WT-10	COMPLETED	Re-transcribe	—	409 Conflict — terminal state

Edge Case Tests

Test	Scenario	Expected
WT-11	Reassignment in ASSIGNED state	Old assignee notified, new assignee set, status stays ASSIGNED
WT-12	Skip from ASSIGNED → REVIEWED	409 Conflict — TRANSCRIBED step cannot be skipped
WT-13	Concurrent state transitions	Only one succeeds; other gets 409 Conflict
WT-14	Reopen COMPLETED transcript	Either 409 (immutable) or creates new workflow instance
WT-15	Timeout in TRANSCRIBED state	Auto-escalation or notification after threshold
WT-16	Review rejection (REVIEWED → TRANSCRIBED)	Status reverts to TRANSCRIBED with rejection notes
WT-17	Audit trail	Every transition logs: timestamp, actor, previous state, new state

Task 5 — Automation Approach

Tools

Category	Tool
Test Framework	pytest (Python) or Jest (Node.js)
API Testing	pytest + requests or Supertest
Assertions	pytest built-in or Chai
CI/CD	GitHub Actions or Jenkins
Linting	flake8 / ESLint
Coverage	pytest-cov / nyc
Mocking	pytest-mock / MSW (Mock Service Worker)

Example Test Structure (pytest)

```
# tests/test_transcript_processing.py

import pytest
import json
import requests
from datetime import datetime

BASE_URL = "http://localhost:8080"
```

```

class TestTranscriptValidation:
    """Task 1 & 2: Transcript quality checks"""

    @pytest.fixture
    def raw_transcript(self):
        with open("transcript_raw_corrupted.json") as f:
            return json.load(f)

    @pytest.fixture
    def metadata(self):
        with open("metadata_corrupted.json") as f:
            return json.load(f)

    def test_speaker_name_consistency(self, metadata):
        """TC-2: first_name + last_name must equal full_name"""
        for speaker in metadata["speakers"]:
            expected = f"{speaker['first_name']} {speaker['last_name']}"
            assert speaker["full_name"] == expected, (
                f"Name mismatch: '{speaker['full_name']}' != '{expected}'"
            )

    def test_word_timestamps_monotonic(self, raw_transcript):
        """TC-3: Word timestamps must be non-decreasing within utterance"""
        for utterance in raw_transcript["utterances"]:
            prev_ts = -1
            for word in utterance["words"]:
                assert word["start"] >= prev_ts, (
                    f"Non-monotonic timestamp: {word['start']} < {prev_ts}"
                )
                prev_ts = word["start"]

    def test_word_within_utterance_bounds(self, raw_transcript):
        """TC-4: Words must not start before their utterance"""
        for utterance in raw_transcript["utterances"]:
            for word in utterance["words"]:
                assert word["start"] >= utterance["start"], (
                    f"Word '{word['text']}' starts at {word['start']} "
                    f"before utterance at {utterance['start']}"
                )

    def test_no_duplicate_word_timestamps(self, raw_transcript):
        """TC-5: No two words share the same timestamp"""
        for utterance in raw_transcript["utterances"]:
            seen = set()
            for word in utterance["words"]:
                assert word["start"] not in seen, (
                    f"Duplicate timestamp {word['start']} for '{word['text']}'"
                )
                seen.add(word["start"])

```

```

def test_created_at_is_valid_iso8601(self, metadata):
    """TC-6: created_at must be valid ISO 8601"""
    dt = datetime.fromisoformat(metadata["created_at"].replace("Z", "+00:00"))
    assert 0 <= dt.hour <= 23
    assert 0 <= dt.minute <= 59

def test_audio_transcript_duration_consistency(self, metadata):
    """TC-7: Transcript must not exceed audio duration"""
    assert metadata["transcript_duration_ms"] <= metadata["audio_duration_ms"]

class TestAPIEndpoints:
    """Task 3: API testing"""

    def test_post_process_valid_payload(self):
        payload = {"utterances": [], "metadata": {}}
        resp = requests.post(f"{BASE_URL}/transcripts/process", json=payload)
        assert resp.status_code in [200, 202]

    def test_post_process_missing_utterances(self):
        payload = {}
        resp = requests.post(f"{BASE_URL}/transcripts/process", json=payload)
        assert resp.status_code == 400

    def test_get_transcript_valid_id(self):
        resp = requests.get(f"{BASE_URL}/transcripts/abc-123")
        assert resp.status_code in [200, 404]

    def test_get_transcript_invalid_id(self):
        resp = requests.get(f"{BASE_URL}/transcripts/; DROP TABLE transcripts;")
        assert resp.status_code in [400, 404]

class TestWorkflowTransitions:
    """Task 4: State machine testing"""

    VALID_TRANSITIONS = [
        ("NEW", "ASSIGNED"),
        ("ASSIGNED", "TRANSCRIBED"),
        ("TRANSCRIBED", "REVIEWED"),
        ("REVIEWED", "COMPLETED"),
    ]

    def test_valid_transitions(self):
        for from_state, to_state in self.VALID_TRANSITIONS:
            resp = requests.post(
                f"{BASE_URL}/workflow/transition",
                json={"state": from_state, "next_state": to_state},
            )
            assert resp.status_code == 200, (
                f"Valid transition {from_state} → {to_state} failed"
            )

```

```
def test_invalid_skip_transitions(self):
    invalid = [
        ("NEW", "COMPLETED"),
        ("NEW", "TRANSCRIBED"),
        ("ASSIGNED", "REVIEWED"),
        ("COMPLETED", "ASSIGNED"),
    ]
    for from_state, to_state in invalid:
        resp = requests.post(
            f"{BASE_URL}/workflow/transition",
            json={"state": from_state, "next_state": to_state},
        )
        assert resp.status_code == 409, (
            f"Invalid transition {from_state} → {to_state} should be rejected"
        )
```

CI/CD Integration (GitHub Actions)

```
# .github/workflows/test.yml
name: QA Tests
on: [push, pull_request]

jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - name: Setup Python
        uses: actions/setup-python@v5
        with:
          python-version: "3.12"
      - name: Install dependencies
        run: pip install pytest requests pytest-cov
      - name: Run transcript quality tests
        run: pytest tests/test_transcript_processing.py -v --cov
      - name: Run API tests (against mock server)
        run: pytest tests/test_api.py -v
      - name: Run workflow tests
        run: pytest tests/test_workflow.py -v
```

Bonus — Scoring System & Programmatic Detection

Transcript Quality Score (0–100)

Category	Weight	Checks
Word Accuracy	30 pts	Levenshtein distance on known entities; flag words with confidence < 0.7

Speaker Attribution	25 pts	All labels mapped to metadata; no unknown speakers; consistent labeling
Timestamp Integrity	20 pts	Monotonic timestamps; no duplicates; words within utterance bounds
Completeness	15 pts	No truncated final utterance; audio/transcript duration match
Formatting	10 pts	No duplicate lines; valid JSON; valid ISO dates

Score = sum of category scores, capped at 100.

Programmatic Detection Rules (Pseudocode)

```
def detect_issues(raw_asr, processed, metadata):
    issues = []

    # Rule 1: Name consistency
    for speaker in metadata["speakers"]:
        if f"{speaker['first_name']} {speaker['last_name']}" !=
speaker["full_name"]:
            issues.append("SPEAKER_NAME_MISMATCH")

    # Rule 2: Timestamp monotonicity
    for utterance in raw_asr["utterances"]:
        for i in range(1, len(utterance["words"])):
            if utterance["words"][i]["start"] < utterance["words"][i-1]["start"]:
                issues.append("NON_MONOTONIC_TIMESTAMP")
            if utterance["words"][i]["start"] == utterance["words"][i-1]["start"]:
                issues.append("DUPLICATE_TIMESTAMP")

    # Rule 3: Speaker label coverage
    transcript_labels = set(u["speaker"] for u in raw_asr["utterances"])
    metadata_labels = set(
        s["post_asr_label"] for s in metadata["speakers"]
    ) | set(
        s["real_time_asr_label"] for s in metadata["speakers"]
    )
    unmapped = transcript_labels - metadata_labels
    for label in unmapped:
        issues.append(f"UNMAPPED_SPEAKER:{label}")

    # Rule 4: Entity name matching
    known_entities = [p for p in metadata.get("parties", [])]
    processed_text = "\n".join(processed.split("\n"))
    for entity in known_entities:
        if entity.lower() not in processed_text.lower():
            issues.append(f"MISSING_ENTITY:{entity}")

    # Rule 5: Date validation
    try:
        dt = datetime.fromisoformat(metadata["created_at"].replace("Z", "+00:00"))
```

```
        if dt.hour > 23 or dt.minute > 59:
            issues.append("INVALID_TIMESTAMP")
    except ValueError:
        issues.append("INVALID_TIMESTAMP")

# Rule 6: Duplicate consecutive lines
lines = processed.strip().split("\n")
for i in range(1, len(lines)):
    if lines[i].strip() == lines[i-1].strip() and lines[i].strip():
        issues.append("DUPLICATE_LINE")

# Rule 7: Truncated ending
last_line = lines[-1].strip()
if last_line and not last_line[-1] in ".!?:":
    issues.append("TRUNCATED_TRANSCRIPT")

return issues
```

Summary

Task	Deliverable	Count
Task 1	Issues identified	20 issues across 4 categories
Task 2	Test cases designed	10 test cases
Task 3	API testing plan	Full plan for both endpoints
Task 4	Workflow tests	17 transition tests (valid, invalid, edge cases)
Task 5	Automation approach	pytest framework + CI/CD pipeline
Bonus	Scoring + detection	100-point rubric + rule-based detector

Playwright Tests Explanation

This document explains the logic inside the `playwright_tests.spec.ts` file, breaking down how the test script works and how it catches the errors present in the provided corrupted test files.

1. Setup & File Loading

```
test.beforeAll(() => {
    rawTranscript = JSON.parse(fs.readFileSync(rawTranscriptPath, 'utf8'));
    metadata = JSON.parse(fs.readFileSync(metadataPath, 'utf8'));
    processedTranscript = fs.readFileSync(processedTranscriptPath, 'utf8');
});
```

What this does: Before any tests run, this code block uses Node.js's built-in `fs` (file system) module to read the corrupted `.json` and `.txt` files from your directory. It parses the JSON files into JavaScript

objects so we can programmatically inspect the data (like checking `metadata.speakers`).

2. Data Integrity Tests (Task 1 & 2)

These tests assert that the data behaves according to logical rules. **Because we are running these tests against intentionally corrupted files, these tests are expected to FAIL.** When they fail, it proves that the testing logic successfully caught the bugs!

TC-1: Entity Names

Checks if the entity names found in `metadata.parties` (e.g., "Cannon Farms") exist perfectly in the processed transcript.

- **Why it fails:** The processed text has AI misspellings like "Cancun Farms" instead of "Cannon Farms".

TC-2: Speaker Role Consistency

Iterates through all speakers in the metadata and checks if `first_name` + `last_name` equals `full_name` .

- **Why it fails:** Speaker A is listed with first name "Terry", but full name "Jerry Sellingman".

TC-3, TC-4, TC-5: Timestamp Integrity

These tests loop through every single word in the raw transcript to verify time logic:

- **Monotonicity (TC-3):** Ensures time always moves forward. (Fails because some words go backward in time).
- **Containment (TC-4):** Ensures a word doesn't start *before* the utterance it belongs to.
- **No Duplicates (TC-5):** Ensures two distinct words don't share the exact same starting millisecond.

TC-6: Valid ISO Date

Takes the `created_at` field from metadata and attempts to parse it into a real JavaScript `Date` object.

- **Why it fails:** The date is `"2025-10-09T26:61:00Z"` . There is no 26th hour or 61st minute, making it an "Invalid Date".

3. API & Workflow Testing (Tasks 3 & 4)

Playwright isn't just for clicking around websites in a browser; it has a powerful backend testing tool called `APIRequestContext` (accessible via the `request` fixture).

```
test('Should return 409 Conflict for invalid skips', async ({ request }) => {
  // const response = await request.post(`${BASE_URL}/workflow/transition`, ...
});
```

What this does: We simulate hitting the endpoints you defined in your `QA_Assessment_Response.md` (`POST /transcripts/process` and `/workflow/transition`).

- We check that sending valid data returns a `200 OK` status.
- We test **Edge Cases**, like trying to skip a workflow state (e.g., going straight from `NEW` to `COMPLETED`), asserting that the API correctly rejects the request by returning a `409 Conflict` status.

(Note: These API calls are mocked/commented out because we don't have an actual backend server running, but they demonstrate exactly how a Software Engineer in Test (SDET) would automate this!)